

PP3-4: Semantic Analysis

Date Due: Checkpoint (pp3): 4/04/16, Final (pp4): 17/04/16

1 Goal

In the third programming project, your job is to implement a semantic analyzer for your compiler. You're now at the penultimate phase of the front-end. If you confirm the source program is free from compile-time errors, you're ready to generate code!

Our semantic analyzer will traverse the parse tree constructed by the parser and validate that the semantic rules of the language are being respected, printing appropriate error messages for violations. This assignment has a bit more room for design decisions than the previous assignments. Your program will be considered correct if it verifies the semantic rules and reports appropriate errors, but there are various ways to go about accomplishing this and ultimately how you structure things is up to you.

Your finished submission will have a working semantic analyzer that reports all varieties of errors. Given that semantic analysis is a bigger job than the scanner or parser, we have plotted a "checkpoint" along the way. At the pp3 checkpoint, your analyzer needs to show it can handle errors related to scoping and declarations, which are the foundation for pp4. We will quickly return feedback, so you will be able to go on to check the rest of the semantic rules.

Past students speak of the semantic analyzer project in hushed and reverent tones. It is a big step up from pp1 and 2. Although an expert procrastinator can start pp1 or pp2 the night before and still crank it out, a one-night shot at pp3 is not recommended. Give yourself plenty of time to work through the issues, post questions to the newsgroup, and thoroughly test your project. In particular, you need to build a solid and robust foundation for your checkpoint to ensure you will be able to complete all the tasks of semantic analysis by the final due date.

2 Semantic Rules of Decaf

Since you are about to embark upon a journey to write a semantic analyzer, you first should know the rules you need to enforce! You will want to carefully read the typing rules, identifier scoping rules, and other restrictions of Decaf as given in the specification handout. Your compiler is responsible for reporting any transgression against the semantic language rules. For each requirement given in the spec, you may want to consider what information you will need to gather to be able to check that requirement and when and where that checking is handled.

3 Error Reporting

Running a semantically correct program through your compiler should not produce any output at this stage. However, when a rule is violated, you are responsible for printing an appropriate message for each error. Your compiler should not stop after the first error, but instead continue checking the rest of the parse tree.

You should use our provided error-reporting facility for printing error messages. `ReportError` will identify the line number and print the source text, underlining the particular section in error. The line that follows is a message describing the problem.

```

*** Error on line 9
here is the offennndnig line with the troubling section underlined
      ~~~~~
Misspelled words are not allowed in Decaf programs!

```

If you remember back to pp1, each token's location was recorded into *yylloc* by the scanner. yacc then tracked the location of each symbol on the parse stack using @1 , @2, etc. As part of pp2, you were storing locations with the parse tree nodes, at the time you may not have realized the eventual purpose was for error-reporting. Those locations are used to provide the context for the error and precisely point out where the trouble is.

Devising clear wording for all the various error situations is actually trickier than it sounds. (Think of all the confusing messages you've seen from compilers... The best compilers have a plethora of specialized error messages, each used in a very particular situation. However, in the interest of keeping things manageable, we will adopt a small set of fairly generic error messages and use each in a variety of contexts. Our pre-defined error messages are listed in *errors.h* and described below. Your output is expected to match our wording and punctuation **exactly**.

4 Error Messages for pp3

For this part, you will report problems with declarations to show us that you have completed the basic functionality for declarations and scoping. The three errors that you are required to catch at the checkpoint are listed below. Those portions in boldface are placeholders, they are replaced with the particulars of the problematic declaration in actual messages.

```

*** Declaration of 'a' here conflicts with declaration on line 5
*** Method 'b' must match inherited type signature

```

These are used to report a new declaration that conflicts with an existing one. The first message is the generic message, used in most contexts (e.g. class redefinition, two parameters of the same name, and so on). The second is a specialized message for a mismatch when a class overrides an inherited method.

```

*** No declaration for class 'Cow' found

```

This message is used to report undeclared identifiers. The only undeclared identifiers you have to catch for the checkpoint are use of undeclared named types in declarations, i.e. a named type used in a variable/function declaration, a class named in an extends clause, or an interface from an implements clause. This error message is also used for undeclared variables and functions, but checking for those is a task handled after the checkpoint.

5 Error messages for pp4

Once you have your foundation built, you will go on to add all the necessary checks to ensure no semantic rules are being violated. Where you find errors, you are to print informative messages to alert the programmer to the root cause. Here is the list of error messages from the full semantic analyzer. As before, the portions in boldface are placeholders.

*** No declaration for function 'Binky' found

Used to report undeclared identifiers (classes, interfaces, functions, variables).

*** Class 'Cow' does not implement entire interface 'Printable'

Used for a class that claims to implement an interface but fails to implement one or more of the required methods.

*** 'this' is only valid within class scope

Used to report use of *this* outside class scope.

*** Incompatible operands: double * string

*** Incompatible operand: ! int[]

Used to report expressions with operands of inappropriate type. Assignment, arithmetic, relational, equality, and logical operators all use the same messages. The first is for binary operators, the second for unary.

*** [] can only be applied to arrays

*** Array subscript must be an integer

*** Size for NewArray must be an integer

Used to report incorrectly formed array subscript expressions or improper use of the *NewArray* built-in.

*** Function 'Winky' expects 4 arguments but 3 given

*** Incompatible argument 2: string given, string[] expected

*** Incompatible argument 3: double given, int/bool/string expected

Used to report mismatches between actual and formal parameters in a function call. The last one is a special-case message used for incorrect argument types to the *Print* built-in function.

*** Cow has no such field 'trunk'

*** Cow field 'spots' only accessible within class scope

Used to report problems with the dot operator. Field means either variable or method. The first message reports to an attempt to apply dot to a non-class type or access a non-existent field from a class. The last is used when the field exists but is not accessible in this context.

*** Test expression must have boolean type

*** break is only allowed inside a loop

*** Incompatible return: int given, void expected

Used to report improper statements.

We have deliberately tried to provide catch-all error messages rather than enumerate distinct messages for all the errors. For example, the “incompatible operands” message is used when the % operator is applied to two strings or when assigning null to a double variable. This will help reduce the number of different errors you need to emit.

6 Error recovery

You will need to determine the appropriate action for your compiler to take after an error. The goal is to report each error once and recover in such a way that few or no further errors will result from the same root cause. For example, if a variable is declared of an undeclared named type, after reporting the error, you might be flexible in the rest of the compilation in allowing that variable to be used. Assume that once the declaration is fixed, it is likely the later uses of it will be correct as well.

There is some guesswork about how to proceed. If you encounter a second declaration that conflicts with the first, do you keep the first and discard the second? Replace the first with the second? Is one strategy more likely to cause fewer cascading errors later? What about if you see a declaration with a mangled type or a completely undeclared variable? Should you try to guess the intended type from the surrounding context? Should you mark the variable as having an error type and use that status to suppress additional errors? How will you constrain errors from tainting the rest of the context, such as adding five operands where the first is a string?

Ideally, you want to continue checking everything else that you can, while suppressing any cascading errors that result from the first error, i.e. those errors that would most likely be fixed if the original error were fixed. A few examples might help. Consider $b[4]$ where b is undeclared. You report about b , but should you also report that you cannot apply brackets to a non-array type? If you assume that if b was supposed to have been declared of the needed array type, the second would also be fixed. What about $b[4] + 5$? Again, if you assume the missing declaration was an int array, this should be fine, so it is another cascading error that should be ignored. What about $b[4.5] = 10$? b is still undeclared, but the array subscript is also not right. Fixing the declaration of b won't fix the array subscript problem, so there are two errors to report. Now consider $b[4.5] = 1 + 4.0$; The error on the right side is yet another distinct error (co-mingled int and double) and fixing the left hand side will not fix that error, so there are three errors to report. The idea is that each distinct error that requires an independent action to correct gets its own error message, but those errors due to the same root cause are not re-reported.

It will come down to judgment calls on the fringe cases. It will be up to you to make decisions and document your reasoning. If you make a design decision that produces output different from the solution/dcc, make sure to document your decision in a README.

7 Starter files

The starting files for this project are in the */pub/projects/cpsc434/assignments/pp3 subdirectory*. The starting project contains the following files (the boldface entries are the ones you will definitely modify, depending on your strategy you may modify others as well):

Makefile	builds project
main.cc	main() and some helper functions
scanner.h/l	our scanner interface/implementation
parser.y	yacc parser for Decaf (replace with your pp2 parser)
ast.h/.cc	interface/implementation of base AST node class
ast_type.h/.cc	interface/implementation of AST type classes
ast_decl.h/.cc	interface/implementation of AST declaration classes
ast_expr.h/.cc	interface/implementation of AST expression classes
ast_stmt.h/.cc	interface/implementation of AST statement classes
errors.h/.cc	error-reporting class for you to use
hashtable.h/.cc	simple hashtable template class
list.h	simple list template class
location.h	utilities for handling locations, yylloc/yyltype
utility.h/.cc	interface/implementation of our provided utility functions
samples/	directory of test input files

Copy the entire directory to your home directory. Use *make* to build the project. The provided makefile will build a compiler called *dcc*. It reads input from stdin, writes output to stdout, and writes errors to stderr. You can use standard UNIX file redirection to read from a file and/or save the output to a file:

```
% dcc <samples/program.decaf >& program.outanderr
```

We provide starter code that will compile but is very incomplete. It will not run properly if given sample files as input, so don't bother trying to run it against these files as given.

As always, the first thing to do is to carefully read all the files we give you and make sure you understand the lay of the land. There is not much different from what you started with for pp2, so it should be fairly quick to hit the ground running. A few notes on the starter files:

- You are given the same parse tree node classes as you started with in pp2. We removed the tree-printing routines to avoid clutter.
- Replace *parser.y* with your pp2 parser. You should not need to make changes to the parser or rearrange the grammar, but you can if you like. You may optionally use the *parser.y* we used in generating the solution. It will be contained in *solution/parser.y*.
- We've added another generic collection class for your use, the *hashtable*. A Hashtable maps string keys to values. Like the List class you've already used, it is templated to allow you to use with all different types of values. This may come in handy for managing symbol tables and scoping information. Read the .h file to learn more about its facilities.

8 Semantic analyzer implementation

For the checkpoint, you need to implement store declarations, manage scopes, and report declaration errors. Here is a quick sketch of the tasks that need to be performed for the milestone:

- Design your strategy for scopes. There are many possible implementations, it's your call to figure out how to proceed. Some questions to get you thinking: What information needs to be recorded with each scope and how will you represent it? What are the different kinds of scopes and do they require any special handling? Where is the scope information stored and how did nodes get access to it? How will you manage entering and exiting scopes? What connections are needed between the levels of nested scopes?
- Once you have a scoping plan, implement it. Our provided hashtable class may come in handy for quick mapping of name to declaration.
- Re-read the Decaf spec about scope visibility – all identifiers in a scope are immediately visible when the scope is entered. Note this is different than C and C++.
- Note there are two separate scopes for a function declaration: one for the parameters and one for the body.
- A class has a scope containing its fields (functions and variables). A subclass inherits all fields of its parent class. There are various ways to handle inheritance (linking the subclass scope to the parent, copying over the fields, etc.). Consider the tradeoffs and choose the strategy you will implement. Interfaces can be handled somewhat similarly to classes.
- Once you have a scoping system in place, when a declaration is entered into scope, you can check for conflicts. And once declarations are stored and can be retrieved, you can verify that all named types used in declarations are valid.
- Add error reporting for conflicting or improper declarations and you're done with the checkpoint. Congrats!

Moving on to the full implementation of the semantic analyzer.

- Start by making sure you are completely familiar with the semantic rules of Decaf as given in the specification handout. Look through the sample files and examine for yourself what are the errors are in the “bad” files and why the good files are well-behaved.
- A design strategy we'd recommend is implementing a polymorphic *Check()* method in the ast classes and do an in-order walk of the tree, visiting and checking each node. Checking on a *VarDecl* might mean ensuring the type is valid and the name is unique for this scope. Checking a *LogicalExpr* could verify both operands are of boolean type. Checking a *BreakStmt* would make sure it is in the context of a loop body.
- Establishing proper behavior for type equivalence and compatibility is a critical step. Re-read the spec and take care with the issues related to inheritance and interfaces. Test this thoroughly in isolation since so much of the expression checking will rely on it working correctly.
- Be sure to take note that many errors are handled similarly (all the arithmetic expressions, for example). With a careful design, you can unify these cases and have less code to write and debug.

- The field access node is one of the trickier parts. Make sure that access to instance variables is properly enforced for the various scopes and that the implicit *this*. Dealing with the *.length()* accessor for arrays will require its own special handling.
- Check out the pseudo base type *errorType*, which can be used to avoid cascading error reports. If you make it compatible with all other types, you can suppress further errors from the underlying problem.
- Testing, testing, and more testing. Make up lots of cases and make sure any fixes you add don't introduce regressions. Before you submit, scan the error messages and semantic rules one last time to make sure you have caught all of them.

9 Matching our output

- Please be sure to take advantage of our provided error-reporting support to make it easy for you to match our output.
- When a file has more than one error, the order the errors are reported is usually correlated to lexical position within the file, i.e. an error on the first line is reported before one on the second and so on. Errors on the same line are usually reported from left to right.
- We will diff your output against the solution, and perform manual checks. If your output varies slightly from the solution (e.g., you print a few errors on the same line in reverse order, you catch the same first error, but the cascading errors are different), please document.

10 Notes on C++

- Feel free to check out the C++ handout in the handouts directory if you're getting lost.
- It will probably come up during your C++ adventures that you will need to downcast some of your node pointers. The ordinary unsafe static C typecast can be used for this purpose, but C++ also has a safe runtime template cast. The syntax is:

```
Decl *d = ...;
FnDecl *fn = dynamic_cast<FnDecl *>(d);
```

This will assign *fn* to the downcast *Decl* if successful or NULL if *d* is not a *FnDecl*.

- Do not feel obligated to worry about memory leaks, even though a real compiler would need to (well, to be honest, even some production compilers leak like a sieve).

11 Testing your semantic analyzer

There are various test files, both correct and incorrect, that are provided for your testing pleasure in the samples directory. As for output, if the source program is correct, there is no output. If the source program has errors, the output consists of an error message for each error.

As always, the provided samples are a limited subset and we will be testing on many more cases. It is part of your job to think about the possibilities and devise your own test cases to ferret out issues not covered by the samples. This is particularly important for the final submission.

Note: the project is focused on semantic errors only. We will not test on syntactically invalid input.

12 Grading

The checkpoint is worth 25 points. We will use this milestone to check your progress, by running your submission through a set of tests and diff'ing your output to our solution. We will report any discrepancies back to you as feedback so you can address these issues before the final submission. The checkpoint deadline is **firm** – in order to get quick turnaround, we cannot accept any late submissions for the checkpoint. We will grade the checkpoint on a strictly pass/fail basis: a submission that works correctly on all or most of our tests will receive full credit (25 points), one that succeeds on at least half our tests will receive about half credit (13 points), otherwise it is a zero. Not hitting the checkpoint is a sad (and imminently avoidable) situation – you will have to do that work for the final submission anyway, and you will be bummed if you've lost your chance to earn 25 free points. Don't let this happen to you!

The final pp4 submission is worth 100 points. We will *diff* your output against our solution as a first pass, expecting that we will likely have made similar choices for most situations, but we will also be examining your output manually. The chief criteria we will use is whether your output makes sense from the point of view of a programmer using your compiler. If your output varies from the provided solution/dcc, please document your design decision in a README.

13 Deliverables

Email your professor a tar.gz of the project directory. Submit the checkpoint as pp3 and the final project as pp4.